

Wine Quality Classification

Neural Network from Scratch with NumPy

Implementation of a Multi-Layer Perceptron

with Backpropagation and AdamW Optimizer

December 2025

The Problem

Dataset: UCI Wine Quality (6,497 samples)

- **11 chemical features → 7 quality classes** (scores 3-9)
- Features: acidity, residual sugar, pH, alcohol, sulfur dioxide, etc.

The Challenge: Extreme Class Imbalance

Quality	3	4	5	6	7	8	9
%	0.5%	3.3%	32.9%	43.7%	16.6%	3.0%	0.1%

77% of wines are quality 5 or 6!

Why Is This Problem Difficult?

1. Extreme class imbalance

- Only 5 samples of quality 9 in entire dataset
- Model tends to ignore rare classes

2. Subjective labels

- Human tasters may disagree by 1-2 points

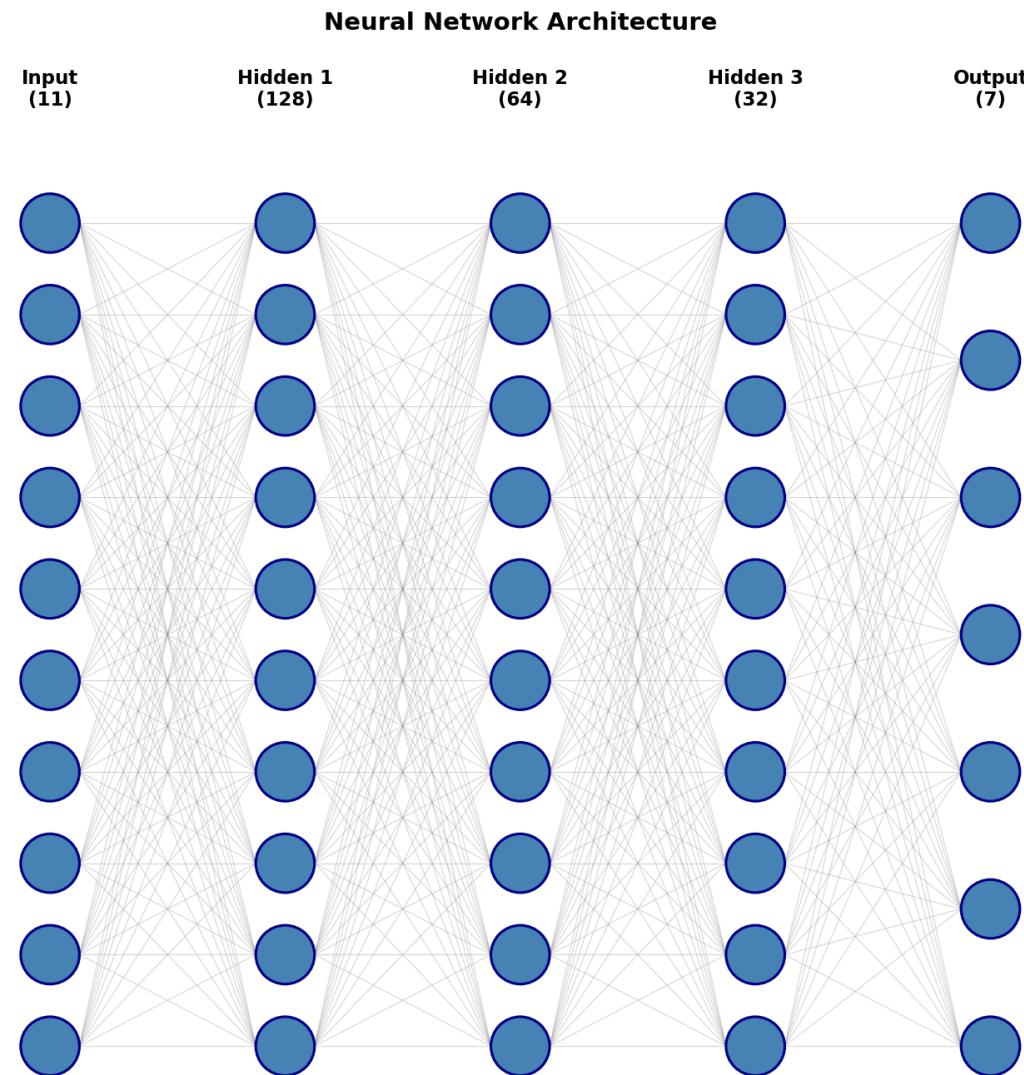
3. Adjacent classes are chemically similar

- Quality 5 and 6 wines have nearly identical compositions

4. Limited features

- No grape variety, aging process, or sensory data

Network Architecture



Why This Architecture?

Layer Design: 128 → 64 → 32

- **First layer (128):** High capacity to learn feature combinations
- **Subsequent layers:** Compress into abstract representations
- **Funnel shape:** Reduces parameters while maintaining power

Why 3 Hidden Layers?

- 1-2 layers: Underfitting
- 3 layers: Best balance ✓
- 4+ layers: No improvement, slower training

Activation Functions

Hidden Layers: Leaky ReLU

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0.01x & \text{if } x \leq 0 \end{cases}$$

Why not standard ReLU?

- ReLU can cause "dying neurons" (gradient = 0 forever)
- Leaky ReLU ensures all neurons remain trainable

Output Layer: Softmax

$$\sigma(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

Weight Initialization

He Initialization

$$W \sim \mathcal{N} \left(0, \sqrt{\frac{2}{n_{in}}} \right)$$

Why He?

- Designed for ReLU-family activations
- Prevents vanishing/exploding signals
- Maintains variance through deep networks

Without proper initialization → unstable training

Loss Function: Cross-Entropy

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N \sum_{c=1}^C y_{i,c} \log(\hat{y}_{i,c})$$

Why Cross-Entropy?

1. Information-theoretic foundation
2. Combined with Softmax \rightarrow clean gradient: $\nabla = \hat{y} - y$
3. Properly penalizes confident wrong predictions

MSE performs poorly for classification tasks

Optimizer: AdamW

AdamW = Adam + Decoupled Weight Decay

Component	Purpose
Momentum ($\beta_1 = 0.9$)	Accelerates convergence
Adaptive LR ($\beta_2 = 0.999$)	Per-parameter learning rates
Weight Decay ($\lambda = 0.001$)	Regularization

Why AdamW over SGD?

- Converged in ~1000 epochs vs ~3000+ with SGD
- Better regularization than standard Adam

Regularization Strategy

The Problem: Severe Overfitting

Without regularization: **77% train, 56% test** (21% gap!)

Technique	Purpose	Setting
Dropout	Prevent co-adaptation	30%
Early Stopping	Stop before memorization	300 epochs patience
LR Decay	Fine-tuning	0.95x every 500 epochs

Result: Gap reduced from 21% → 4%

Why 30% Dropout?

Tested multiple dropout rates:

Rate	Result
0-20%	Still significant overfitting
30%	Best balance ✓
40-50%	Underfitting, lower accuracy

How it works:

- Randomly zero out 30% of neurons during training
- Forces network to learn robust features
- Acts like training ensemble of networks

Data Preprocessing

Z-score Normalization

$$x' = \frac{x - \mu}{\sigma}$$

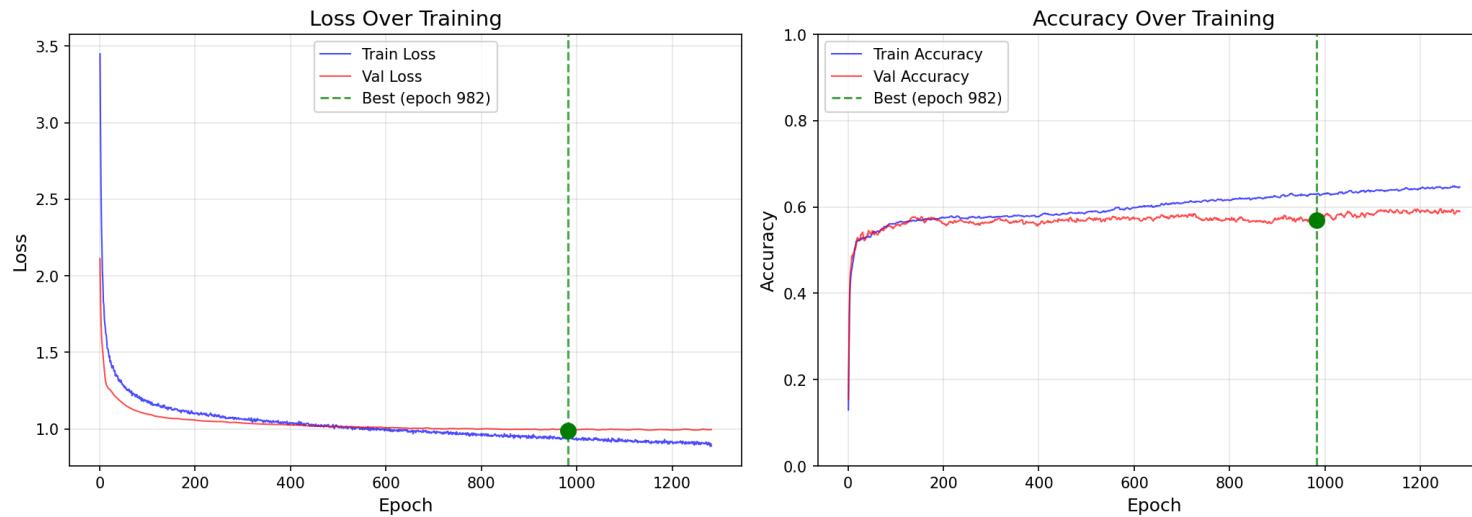
Features have different scales:

- Density: 0.99 - 1.04
- Sulfur dioxide: 6 - 440

Stratified Split (70% / 15% / 15%)

- Maintains class proportions in each split
- Critical for imbalanced data

Training Results



Early stopping at epoch ~1282, restored best weights from epoch 982

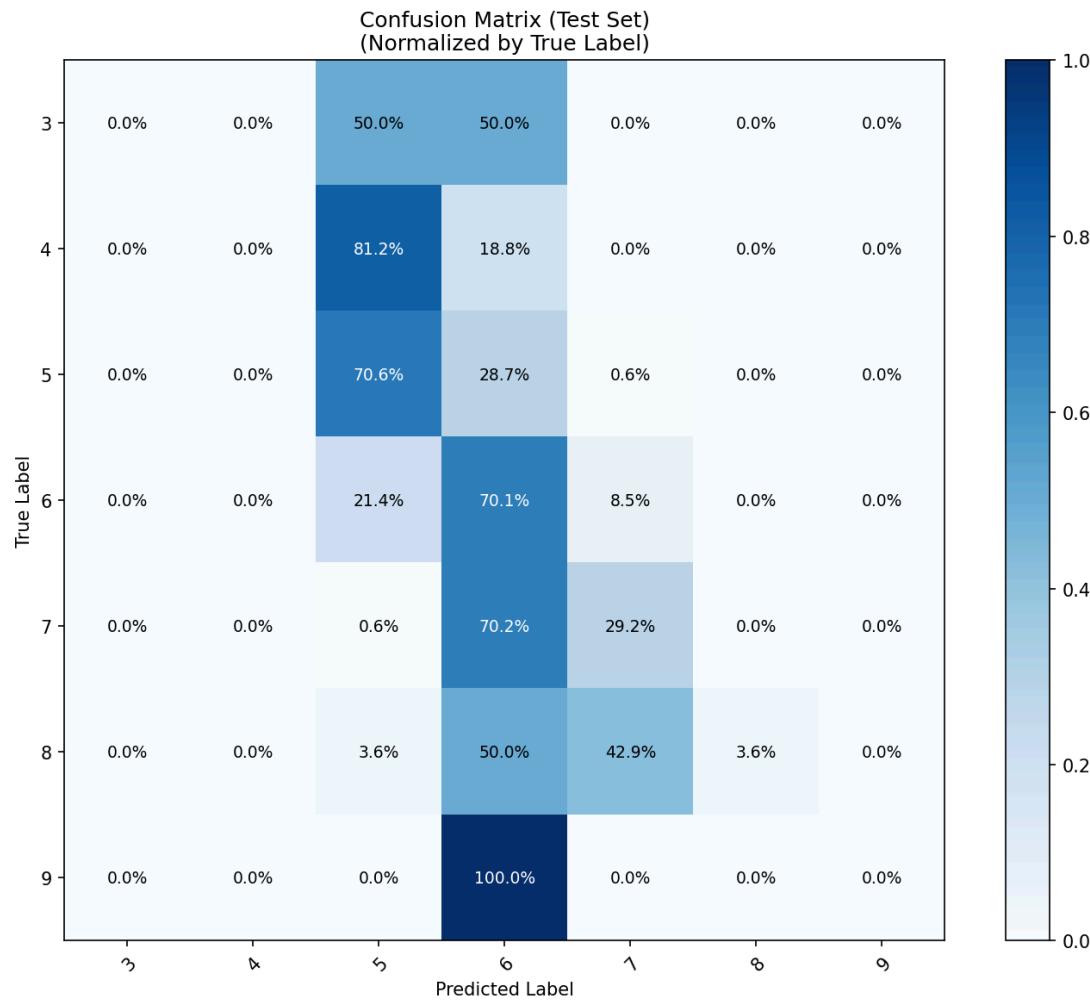
Final Performance

Metric	Value
Training Accuracy	63.03%
Validation Accuracy	56.95%
Test Accuracy	58.91%

Comparison with Baselines

Method	Accuracy
Random guessing	14.3%
Always predict "6"	43.7%
Our model	58.9%

Confusion Matrix Analysis



Error Analysis

What the model does well:

- Classes 5, 6, 7 (sufficient training data)
- ~65% recall on majority classes

What the model struggles with:

- Classes 3 and 9 (almost never predicted)
- Adjacent class confusion (5



6, 6



7)

Why ~59% Accuracy is Good

1. **vs Random:** 14.3% → 58.9% = **4.1x improvement**
2. **vs Majority baseline:** 43.7% → 58.9% = **+15 points**
3. **vs Literature:** Published results with SVM, Random Forest, XGBoost achieve 55-65%
4. **Problem difficulty:** Even human experts disagree on adjacent quality levels

Our from-scratch implementation is competitive!

Limitations

1. Class Imbalance

- Model ignores extremely rare classes (3, 9)
- Would need class weighting or oversampling

2. Feature Limitations

- Only 11 chemical measurements
- Missing: grape variety, vintage, winemaking process

3. Ordinal vs Categorical

- We treat quality as 7 independent classes
- Being wrong by 1 should cost less than wrong by 5

Potential Improvements

Technique	How it helps
Class weighting	Penalize minority class errors more
SMOTE	Generate synthetic samples for rare classes
Ordinal regression	Respect the natural order of quality
Ensemble methods	Combine multiple models
Feature engineering	Polynomial features, interactions

Conclusion

What we built:

- Neural network **100% from scratch** with NumPy
- Forward/backward propagation, AdamW, dropout — all manual

Key results:

- **~59% test accuracy** on 7-class imbalanced problem
- Significantly outperforms baselines
- Competitive with literature results

Questions?

Repository

github.com/fraco-oxza/final-ia

Key files:

- `neural_network.py` — Full implementation
- `REPORT.md` — Detailed documentation